

```
# Importamos las librerias para el analisis exploratorio
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
sns.set_style('whitegrid')
plt.style.use("fivethirtyeight")
%matplotlib inline
```

```
# Cargamos los datasets
```

```
train =
pd.read_csv("/content/drive/MyDrive/CaixaBank-Hackaton/train.csv")
test =
pd.read_csv("/content/drive/MyDrive/CaixaBank-Hackaton/test_x.csv")
```

```
# Mostramos los 5 primeros registros del dataset de entrenamiento
```

```
train.head(5)
```

	Date	Open	High	Low	Close \
0	1994-01-03	3615.199951	3654.699951	3581.000000	3654.500000
1	1994-01-04	3654.500000	3675.500000	3625.100098	3630.300049
2	1994-01-05	3625.199951	3625.199951	3583.399902	3621.199951
3	1994-01-06	NaN	NaN	NaN	NaN
4	1994-01-07	3621.199951	3644.399902	3598.699951	3636.399902

	Adj Close	Volume	Target
0	3654.496338	0.0	0
1	3630.296387	0.0	1
2	3621.196289	0.0	1
3	NaN	NaN	0
4	3636.396240	0.0	1

```
# Vemos informacion estadistica al detalle de los datos de entrenamiento
```

```
train.describe(include='all')
```

	Date	Open	High	Low
Close \				
count	6554	6421.000000	6421.000000	6421.000000
6421.000000				
unique	6554	NaN	NaN	NaN
NaN				
top	1994-01-03	NaN	NaN	NaN
NaN				
freq	1	NaN	NaN	NaN
NaN				

mean	NaN	8936.540448	9005.235576	8858.340567
8934.978558				
std	NaN	2732.102441	2749.009324	2712.511028
2731.032625				
min	NaN	2865.100098	2877.300049	2833.600098
2865.100098				
25%	NaN	7732.399902	7817.200195	7641.500000
7727.799805				
50%	NaN	9329.700195	9404.599609	9243.000000
9331.000000				
75%	NaN	10525.500000	10590.299805	10441.200195
10523.400391				
max	NaN	15999.200195	16040.400391	15868.599609
15945.700195				

	Adj Close	Volume	Target
count	6421.000000	6.421000e+03	6554.000000
unique	NaN	NaN	NaN
top	NaN	NaN	NaN
freq	NaN	NaN	NaN
mean	8934.970624	8.218074e+07	0.516936
std	2731.030170	1.231845e+08	0.499751
min	2865.097168	0.000000e+00	0.000000
25%	7727.791992	0.000000e+00	0.000000
50%	9331.000000	1.966000e+05	1.000000
75%	10523.400391	1.773980e+08	1.000000
max	15945.683594	7.894902e+08	1.000000

Tenemos datos nulos y hay que buscar un metodo de tratar esos datos

```
train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6554 entries, 0 to 6553
Data columns (total 8 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Date        6554 non-null   object
 1   Open        6421 non-null   float64
 2   High        6421 non-null   float64
 3   Low         6421 non-null   float64
 4   Close       6421 non-null   float64
 5   Adj Close   6421 non-null   float64
 6   Volume      6421 non-null   float64
 7   Target      6554 non-null   int64
dtypes: float64(6), int64(1), object(1)
memory usage: 409.8+ KB
```

Convierto la columna Date a formato fecha
train['Date'] = pd.to_datetime(train['Date'])

```
# Relleno los datos faltantes mediante una interpolacion lineal
```

```
train_fixed = train.copy(deep=True)
columns_to_interpolated = ['Open', 'High', 'Low', 'Close', 'Adj
Close', 'Volume']
train_fixed[columns_to_interpolated] =
train_fixed[columns_to_interpolated].interpolate(method='linear',
axis=0)
train_fixed.head(5)
```

	Date	Open	High	Low	Close
Adj Close \					
0	1994-01-03	3615.199951	3654.699951	3581.000000	3654.500000
					3654.496338
1	1994-01-04	3654.500000	3675.500000	3625.100098	3630.300049
					3630.296387
2	1994-01-05	3625.199951	3625.199951	3583.399902	3621.199951
					3621.196289
3	1994-01-06	3623.199951	3634.799926	3591.049926	3628.799926
					3628.796264
4	1994-01-07	3621.199951	3644.399902	3598.699951	3636.399902
					3636.396240


	Volume	Target
0	0.0	0
1	0.0	1
2	0.0	1
3	0.0	0
4	0.0	1

```
# Ahora no tenemos datos NaN en el dataset
```

```
train_fixed.info()
```

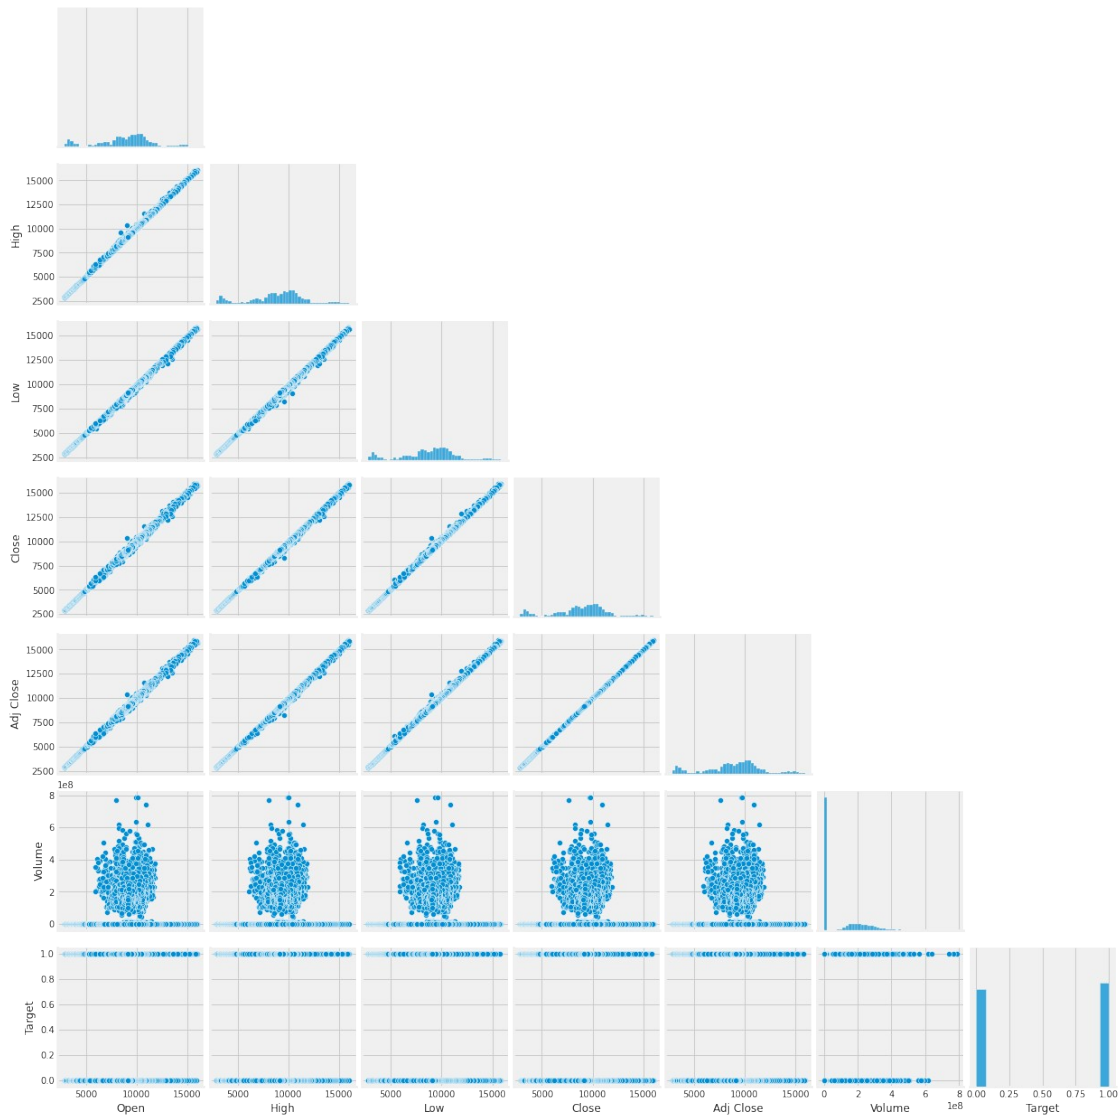
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6554 entries, 0 to 6553
Data columns (total 8 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Date        6554 non-null   datetime64[ns]
1   Open        6554 non-null   float64
2   High        6554 non-null   float64
3   Low         6554 non-null   float64
4   Close       6554 non-null   float64
5   Adj Close   6554 non-null   float64
6   Volume      6554 non-null   float64
7   Target      6554 non-null   int64
dtypes: datetime64[ns](1), float64(6), int64(1)
memory usage: 409.8 KB
```

```
plt.figure(figsize=(12,8))
sns.heatmap(train_fixed.drop(columns=['Target']).corr(), annot=True,
cmap='crest')
plt.title('Matriz de correlacion', fontsize=26)
Text(0.5, 1.0, 'Matriz de correlacion')
```



```
plt.figure(figsize=(2,2))
sns.pairplot(train_fixed, corner=True)
plt.show()
```

<Figure size 144x144 with 0 Axes>



```

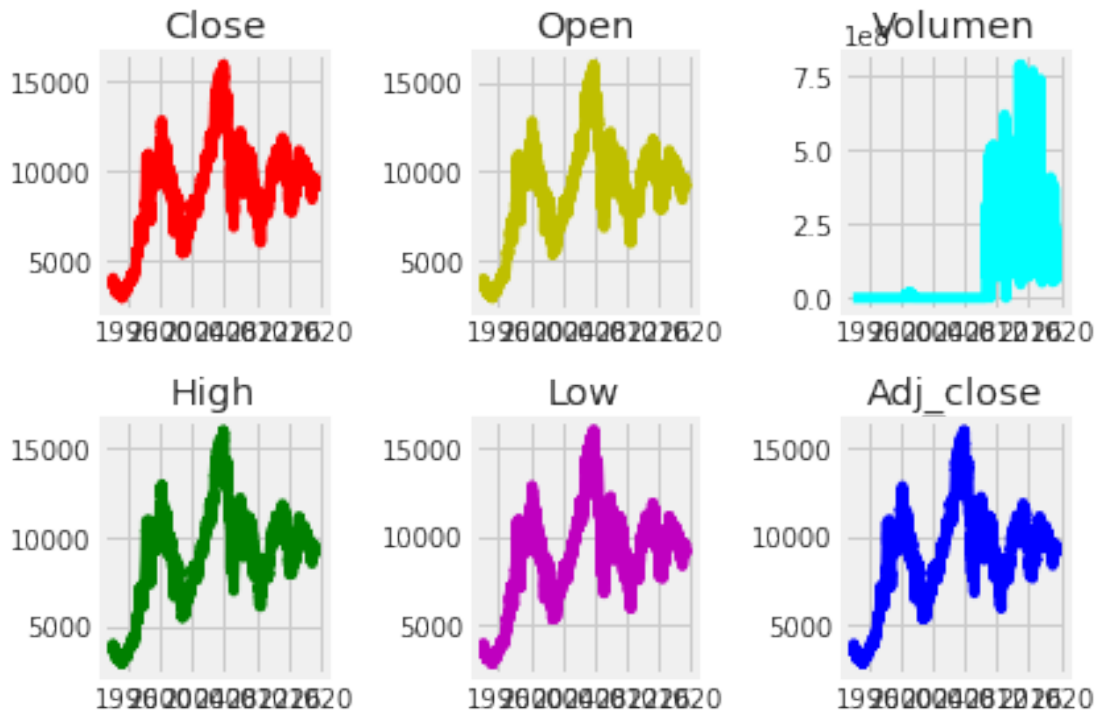
date = train_fixed['Date']
close = train_fixed['Close']
Open = train_fixed['Open']
High = train_fixed['High']
Low = train_fixed['Close']
Adj_close = train_fixed['Adj Close']
Volumen = train_fixed['Volume']
fig , ((axes1,axes2,axes3),(axes4,axes5,axes6))= plt.subplots(2,3)
axes1.plot(date,close,'r')
axes1.set_title('Close')
axes2.plot(date,Open,'y')
axes2.set_title('Open')
axes3.plot(date,Volumen,'cyan')
axes3.set_title('Volumen')
axes4.plot(date,High,'g')
axes4.set_title('High')
axes5.plot(date,Low,'m')

```

```

axes5.plot(date,Low,'m')
axes5.set_title('Low')
axes6.plot(date,Adj_close,'b')
axes6.set_title('Adj_close')
fig.tight_layout()
plt.show()

```



De esta matriz de correlacion notamos que la variable volumen está pobremente correlacionada o directamente no hay relación. **Ya que las otras variables están directamente relacionadas nos brindan la misma informacion y podemos usar solo una para hacer nuestras predicciones, en este caso tomaré la variable 'Close' para hacer las predicciones.**

Vista historica del precio de cierre

```

plt.figure(figsize=(16,8))
plt.title('Close Price History')
plt.plot(train_fixed['Close'])
plt.xlabel('Date', fontsize=18)
plt.ylabel('Close Price USD ($)', fontsize=18)
plt.show()

```



```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
# TensorFlow y tf.keras
import tensorflow as tf
from tensorflow import keras

print(tf.__version__)

2.8.0

X, y = train_fixed.Close.values.reshape(-1, 1),
train_fixed.Target.values.reshape(-1, 1)

# Escalamos nuestros datos en el rango de 0 a 1

escaler = MinMaxScaler()

X = escaler.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.25, shuffle=False)

X_valid = X_test[:X_test.size//2]
X_test = X_test[X_test.size//2:]
y_valid = y_test[:y_test.size//2]
y_test = y_test[y_test.size//2:]

print(f'x_train.shape = {X_train.shape}')
print(f'y_train.shape = {y_train.shape}')
print(f'X_valid.shape = {X_valid.shape}')
print(f'y_valid.shape = {y_valid.shape}')
```

```

print(f'x_test.shape = {X_test.shape}')
print(f'y_test.shape = {y_test.shape}')

x_train.shape = (4915, 1)
y_train.shape = (4915, 1)
X_valid.shape = (819, 1)
y_valid.shape = (819, 1)
x_test.shape = (820, 1)
y_test.shape = (820, 1)

from tensorflow.keras import regularizers

model_basic = keras.Sequential([
    keras.layers.LSTM(
        units=32,
        activation='tanh',
        return_sequences=True,
        recurrent_activation='sigmoid',
        use_bias=True, dropout=0.2,
        kernel_regularizer = regularizers.l2(1e-5),
        input_shape=(X_train.shape[1], 1)),
    keras.layers.Dense(64, activation='relu'),
    keras.layers.LSTM(units=128, activation='tanh',
return_sequences=False, recurrent_activation='sigmoid', use_bias=True,
dropout=0.2, kernel_regularizer = regularizers.l2(1e-5)),
    keras.layers.Dense(32, activation='relu'),
    keras.layers.Dense(1, activation='sigmoid')
])

model_basic.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
lstm (LSTM)	(None, 1, 32)	4352
dense (Dense)	(None, 1, 64)	2112
lstm_1 (LSTM)	(None, 128)	98816
dense_1 (Dense)	(None, 32)	4128
dense_2 (Dense)	(None, 1)	33
=====		
Total params: 109,441		
Trainable params: 109,441		
Non-trainable params: 0		


```
model_basic.compile(  
    metrics=['accuracy'],  
    optimizer='adam',  
    loss='binary_crossentropy'  
)
```

```
history = model_basic.fit(X_train, y_train, epochs=20,  
validation_data=(X_valid, y_valid))
```

Epoch 1/20

154/154 [=====] - 12s 23ms/step - loss: 0.6935 - accuracy: 0.5009 - val_loss: 0.6917 - val_accuracy: 0.5409

Epoch 2/20

154/154 [=====] - 2s 13ms/step - loss: 0.6930 - accuracy: 0.5123 - val_loss: 0.6917 - val_accuracy: 0.5409

Epoch 3/20

154/154 [=====] - 2s 12ms/step - loss: 0.6930 - accuracy: 0.5123 - val_loss: 0.6911 - val_accuracy: 0.5409

Epoch 4/20

154/154 [=====] - 2s 11ms/step - loss: 0.6929 - accuracy: 0.5123 - val_loss: 0.6911 - val_accuracy: 0.5409

Epoch 5/20

154/154 [=====] - 2s 12ms/step - loss: 0.6930 - accuracy: 0.5123 - val_loss: 0.6911 - val_accuracy: 0.5409

Epoch 6/20

154/154 [=====] - 2s 10ms/step - loss: 0.6929 - accuracy: 0.5123 - val_loss: 0.6914 - val_accuracy: 0.5409

Epoch 7/20

154/154 [=====] - 2s 11ms/step - loss: 0.6929 - accuracy: 0.5123 - val_loss: 0.6916 - val_accuracy: 0.5409

Epoch 8/20

154/154 [=====] - 1s 8ms/step - loss: 0.6929 - accuracy: 0.5123 - val_loss: 0.6914 - val_accuracy: 0.5409

Epoch 9/20

154/154 [=====] - 1s 7ms/step - loss: 0.6929 - accuracy: 0.5123 - val_loss: 0.6917 - val_accuracy: 0.5409

Epoch 10/20

154/154 [=====] - 1s 7ms/step - loss: 0.6930 - accuracy: 0.5123 - val_loss: 0.6914 - val_accuracy: 0.5409

Epoch 11/20

154/154 [=====] - 1s 7ms/step - loss: 0.6929 - accuracy: 0.5123 - val_loss: 0.6914 - val_accuracy: 0.5409

Epoch 12/20

154/154 [=====] - 1s 7ms/step - loss: 0.6929 - accuracy: 0.5123 - val_loss: 0.6912 - val_accuracy: 0.5409

Epoch 13/20

154/154 [=====] - 1s 7ms/step - loss: 0.6929 - accuracy: 0.5123 - val_loss: 0.6915 - val_accuracy: 0.5409

Epoch 14/20

154/154 [=====] - 1s 7ms/step - loss: 0.6929 - accuracy: 0.5123 - val_loss: 0.6916 - val_accuracy: 0.5409

```

Epoch 15/20
154/154 [=====] - 1s 6ms/step - loss: 0.6929
- accuracy: 0.5123 - val_loss: 0.6916 - val_accuracy: 0.5409
Epoch 16/20
154/154 [=====] - 1s 6ms/step - loss: 0.6929
- accuracy: 0.5123 - val_loss: 0.6918 - val_accuracy: 0.5409
Epoch 17/20
154/154 [=====] - 1s 7ms/step - loss: 0.6929
- accuracy: 0.5123 - val_loss: 0.6914 - val_accuracy: 0.5409
Epoch 18/20
154/154 [=====] - 1s 7ms/step - loss: 0.6929
- accuracy: 0.5123 - val_loss: 0.6913 - val_accuracy: 0.5409
Epoch 19/20
154/154 [=====] - 1s 7ms/step - loss: 0.6929
- accuracy: 0.5123 - val_loss: 0.6914 - val_accuracy: 0.5409
Epoch 20/20
154/154 [=====] - 1s 7ms/step - loss: 0.6929
- accuracy: 0.5123 - val_loss: 0.6913 - val_accuracy: 0.5409

results = model_basic.evaluate(X_test, y_test)

26/26 [=====] - 0s 4ms/step - loss: 0.6924 -
accuracy: 0.5207

```

```

def visualizacion_resultados(history):
    epochs = [i for i in range(20)]
    fig, ax = plt.subplots(1,2)
    train_acc = history.history["accuracy"]
    train_loss = history.history["loss"]
    val_acc = history.history["val_accuracy"]
    val_loss = history.history["val_loss"]
    fig.set_size_inches(16, 9)

    ax[0].plot(epochs, train_acc, "go-", label = " Entrenamiento
accuracy")
    ax[0].plot(epochs, val_acc, "ro-", label= "Validacion accuracy")
    ax[0].set_title("Entrenamiento & validación accuracy")
    ax[0].legend()
    ax[0].set_xlabel("Epochs")
    ax[0].set_ylabel("Accuracy")

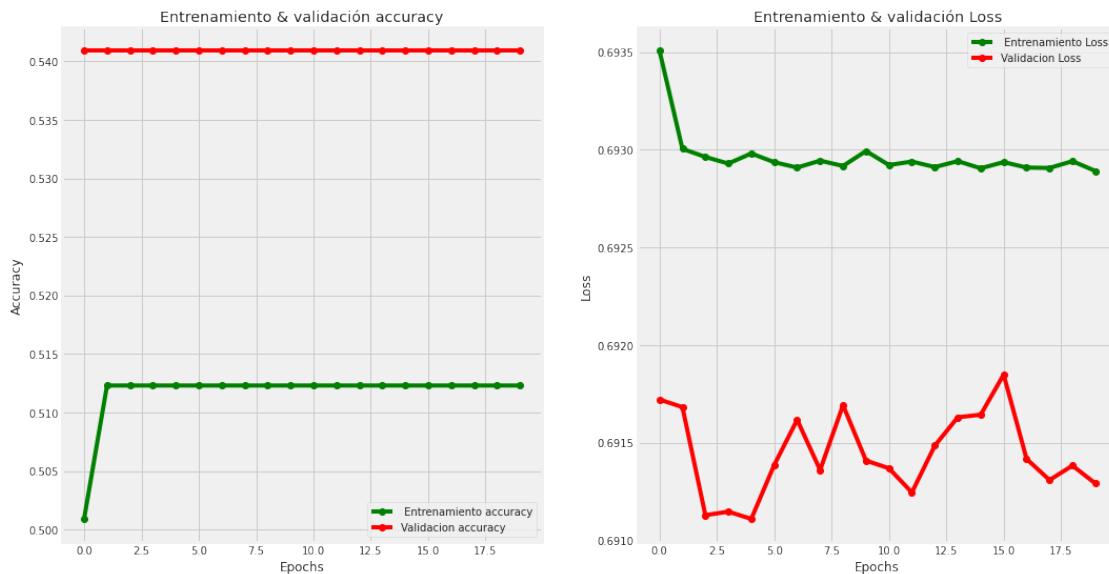
    ax[1].plot(epochs, train_loss, "go-", label = " Entrenamiento Loss")
    ax[1].plot(epochs, val_loss, "ro-", label= "Validacion Loss")
    ax[1].set_title("Entrenamiento & validación Loss")
    ax[1].legend()
    ax[1].set_xlabel("Epochs")
    ax[1].set_ylabel("Loss")

    plt.show()

print(history.history.keys())

```

```
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
visualizacion_resultados(history)
```



```
model_optimizer = keras.Sequential([
    keras.layers.LSTM(
        units=64,
        activation='tanh',
        return_sequences=True,
        recurrent_activation='sigmoid',
        use_bias=True, dropout=0.2,
        kernel_regularizer = regularizers.l2(1e-5),
        input_shape=(X_train.shape[1], 1)),
    keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(256, kernel_regularizer =
regularizers.l2(1e-5), activation = "relu"),
    keras.layers.LSTM(units=128, activation='tanh',
return_sequences=False, recurrent_activation='sigmoid', use_bias=True,
dropout=0.2, kernel_regularizer = regularizers.l2(1e-5)),
    keras.layers.Dense(64, activation='relu'),
    keras.layers.Dense(16, activation='relu'),
    keras.layers.Dense(1, activation='sigmoid')
])
```

```
model_optimizer.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
lstm_2 (LSTM)	(None, 1, 64)	16896
dense_3 (Dense)	(None, 1, 64)	4160

dense_4 (Dense)	(None, 1, 256)	16640
lstm_3 (LSTM)	(None, 128)	197120
dense_5 (Dense)	(None, 64)	8256
dense_6 (Dense)	(None, 16)	1040
dense_7 (Dense)	(None, 1)	17

```
=====
Total params: 244,129
Trainable params: 244,129
Non-trainable params: 0
=====
```

```
model_optimizer.compile(
    metrics=['accuracy'],
    optimizer='adam',
    loss='binary_crossentropy'
)
```

```
history_optimizer = model_optimizer.fit(X_train, y_train, epochs=20,
validation_data=(X_valid, y_valid))
```


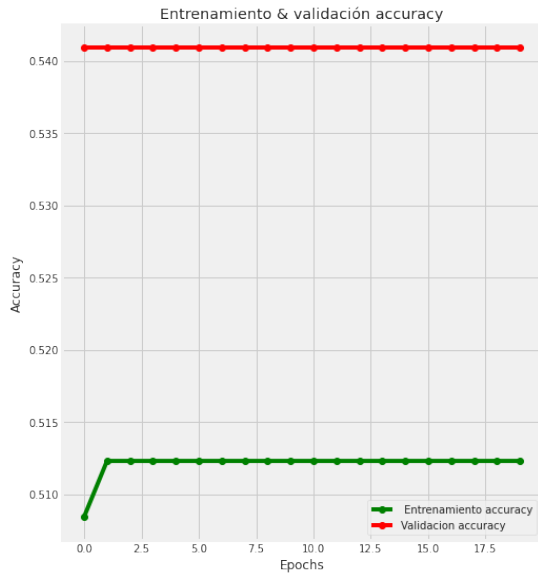
```
Epoch 1/20
154/154 [=====] - 5s 13ms/step - loss: 0.6945
- accuracy: 0.5084 - val_loss: 0.6918 - val_accuracy: 0.5409
Epoch 2/20
154/154 [=====] - 1s 7ms/step - loss: 0.6931
- accuracy: 0.5123 - val_loss: 0.6914 - val_accuracy: 0.5409
Epoch 3/20
154/154 [=====] - 1s 7ms/step - loss: 0.6930
- accuracy: 0.5123 - val_loss: 0.6916 - val_accuracy: 0.5409
Epoch 4/20
154/154 [=====] - 1s 7ms/step - loss: 0.6929
- accuracy: 0.5123 - val_loss: 0.6917 - val_accuracy: 0.5409
Epoch 5/20
154/154 [=====] - 1s 7ms/step - loss: 0.6929
- accuracy: 0.5123 - val_loss: 0.6916 - val_accuracy: 0.5409
Epoch 6/20
154/154 [=====] - 1s 8ms/step - loss: 0.6929
- accuracy: 0.5123 - val_loss: 0.6916 - val_accuracy: 0.5409
Epoch 7/20
154/154 [=====] - 1s 8ms/step - loss: 0.6929
- accuracy: 0.5123 - val_loss: 0.6917 - val_accuracy: 0.5409
Epoch 8/20
154/154 [=====] - 1s 8ms/step - loss: 0.6929
- accuracy: 0.5123 - val_loss: 0.6915 - val_accuracy: 0.5409
Epoch 9/20
```

```
154/154 [=====] - 1s 8ms/step - loss: 0.6929
- accuracy: 0.5123 - val_loss: 0.6915 - val_accuracy: 0.5409
Epoch 10/20
154/154 [=====] - 1s 7ms/step - loss: 0.6929
- accuracy: 0.5123 - val_loss: 0.6915 - val_accuracy: 0.5409
Epoch 11/20
154/154 [=====] - 1s 8ms/step - loss: 0.6929
- accuracy: 0.5123 - val_loss: 0.6916 - val_accuracy: 0.5409
Epoch 12/20
154/154 [=====] - 1s 8ms/step - loss: 0.6929
- accuracy: 0.5123 - val_loss: 0.6913 - val_accuracy: 0.5409
Epoch 13/20
154/154 [=====] - 1s 7ms/step - loss: 0.6929
- accuracy: 0.5123 - val_loss: 0.6915 - val_accuracy: 0.5409
Epoch 14/20
154/154 [=====] - 1s 7ms/step - loss: 0.6929
- accuracy: 0.5123 - val_loss: 0.6914 - val_accuracy: 0.5409
Epoch 15/20
154/154 [=====] - 1s 8ms/step - loss: 0.6929
- accuracy: 0.5123 - val_loss: 0.6916 - val_accuracy: 0.5409
Epoch 16/20
154/154 [=====] - 1s 7ms/step - loss: 0.6930
- accuracy: 0.5123 - val_loss: 0.6916 - val_accuracy: 0.5409
Epoch 17/20
154/154 [=====] - 1s 8ms/step - loss: 0.6929
- accuracy: 0.5123 - val_loss: 0.6915 - val_accuracy: 0.5409
Epoch 18/20
154/154 [=====] - 1s 7ms/step - loss: 0.6929
- accuracy: 0.5123 - val_loss: 0.6914 - val_accuracy: 0.5409
Epoch 19/20
154/154 [=====] - 1s 8ms/step - loss: 0.6929
- accuracy: 0.5123 - val_loss: 0.6917 - val_accuracy: 0.5409
Epoch 20/20
154/154 [=====] - 1s 8ms/step - loss: 0.6929
- accuracy: 0.5123 - val_loss: 0.6914 - val_accuracy: 0.5409

results = model_basic.evaluate(X_test, y_test)

26/26 [=====] - 0s 4ms/step - loss: 0.6924 -
accuracy: 0.5207

visualizacion_resultados(history_optimizer)
```



Predicciones del dataset test.csv

```
test.head()
```

	test_index	Date	Open	High	Low
Close \					
0	6557	2019-06-05	9136.799805	9173.400391	9095.000000
					9150.500000
1	6558	2019-06-06	9169.200195	9246.200195	9136.700195
					9169.200195
2	6559	2019-06-07	9186.700195	9261.400391	9185.700195
					9236.099609
3	6560	2019-06-10	9284.200195	9302.200195	9248.099609
					9294.099609
4	6561	2019-06-11	9288.599609	9332.500000	9273.400391
					9282.099609

	Adj Close	Volume
0	9150.500000	158753000.0
1	9169.200195	212720900.0
2	9236.099609	150664700.0
3	9294.099609	102323700.0
4	9282.099609	144701200.0

Hacemos la prediccion

```
input_test = test.Close.values.reshape(-1,1)
input_test = escaler.fit_transform(input_test)
predictions_test = model_optimizer.predict(input_test, batch_size=64)
predictions_test
predictions_test.reshape(-1,1)
```

```

predictions_test = pd.Series(predictions_test.reshape(-1,),
name='Target')

predictions_test = predictions_test.apply(lambda x: 1 if x >= 0.5 else
0)

predictions_test

prediction = pd.concat([test['test_index'],predictions_test],axis=1)
prediction

   test_index  Target
0         6557       1
1         6558       1
2         6559       1
3         6560       1
4         6561       1
..         ...     ...
721        7278       1
722        7279       1
723        7280       1
724        7281       1
725        7282       1

[726 rows x 2 columns]

prediction.to_csv('predictions.csv')
prediction.to_json('predictions.json')

```